

In [ ]:

```
In [2]: import matplotlib.pyplot as plt
%matplotlib inline

import numpy as np
import seaborn as sns
import pandas as pd
pd.set_option("display.max_columns", 200)

import warnings
warnings.filterwarnings('ignore')
```

## Read datasets from CSV

```
In [3]: df = pd.read_csv("ids_data/attack_dataset.csv.gz") # attack dataset
bonafide = pd.read_csv('ids_data/bonafide_dataset_20191121.csv.gz') # bonafide t
bonafide = pd.concat([bonafide, pd.read_csv('ids_data/bonafide_dataset_20201110.
bonafide = pd.concat([bonafide, pd.read_csv('ids_data/bonafide_dataset_20201129.
```

```
In [4]: # add label to bonafide dataset
bonafide['label'] = "bonafide"
```

## compare both attack and bonafide datasets

```
In [5]: # checking if both datasets have same columns count
if df.shape[1]==bonafide.shape[1]:
    print("Both datasets have the same number of columns")

print(df.shape)
print(bonafide.shape)
```

```
Both datasets have the same number of columns
(455503, 42)
(380438, 42)
```

```
In [6]: # total number of records including attack and bonafide
total = df.shape[0] + bonafide.shape[0]
print(f"total number of records: {total}")

# total number of attack records
malicious = (df.shape[0]/total)*100
print("total % of attack records: {:.2f}".format(malicious))

# total number of bonafide records
legitimate = (bonafide.shape[0]/total)*100
print("total % of bonafide records: {:.2f}".format(legitimate))
```

```
total number of records: 835941
total % of attack records: 54.49
total % of bonafide records: 45.51
```

## bonafide and attack dataset overview

```
In [7]: df.head(3)
```

```
Out[7]:
```

	frame_info.encap_type	frame_info.time	frame_info.time_epoch	frame_info.number
0	1	Sep 2, 2020 21:04:37.063530000 -03	1.599091e+09	1
1	1	Sep 2, 2020 21:04:39.363792000 -03	1.599091e+09	2
2	1	Nov 16, 2020 18:15:14.851050000 -03	1.605561e+09	1

```
In [8]: df.label.value_counts()
```

```
Out[8]: label
zmap          74613
nmap_connect  45882
hping_syn     43750
unicorn_syn   43039
nmap_syn      40642
unicorn_conn  39170
masscan       21138
nmap_ack      20497
nmap_window   18851
nmap_null     12511
nmap_xmas     12505
nmap_fin      12504
hping_ack     11344
nmap_maimon   10493
unicorn_ack    8494
hping_xmas    7344
hping_null    7344
hping_fin     7344
unicorn_null   4690
unicorn_xmas   4466
unicorn_fxmas  4444
unicorn_fin    4438
Name: count, dtype: int64
```

```
In [9]: bonafide.head(3)
```

```
Out[9]:
```

	frame_info.encap_type	frame_info.time	frame_info.time_epoch	frame_info.number
0	1	Nov 21, 2019 02:00:00.309420000 -03	1.574312e+09	2
1	1	Nov 21, 2019 02:00:00.313671000 -03	1.574312e+09	7
2	1	Nov 21, 2019 02:00:00.315642000 -03	1.574312e+09	10

```
In [10]: bonafide.label.value_counts()
```

```
Out[10]: label  
bonafide    380438  
Name: count, dtype: int64
```

## combine both attack and bonafide datasets

```
In [11]: full_data = pd.concat([bonafide, df])
```

---

## Pre-processing

- converted features from hexadecimal to integer

- imputed null value fields by 0

```
In [12]: hex_fields = ['eth.type', 'ip.id', 'ip.flags', 'ip.checksum', 'ip.dsfield', 'tcp  
  
full_data = full_data.fillna(0)  
  
for field in hex_fields:  
    full_data[field] = full_data[field].apply(lambda x: int(str(x), 16))
```

```
In [13]: print(full_data.shape)
```

```
(835941, 42)
```

## check only packets with tcp protocol ( ip proto 6) exist and filter other packets

```
In [14]: non_tcp_records = full_data[full_data['ip.proto'] != 6].shape[0]  
print("Removed", non_tcp_records, "non-tcp packets from the original dataset.")  
  
# includes only the packets/records that are tcp-based  
full_data = full_data[full_data['ip.proto'] == 6]
```

Removed 52177 non-tcp packets from the original dataset.

```
In [15]: full_data.label.value_counts()
```

```
Out[15]: label
bonafide      328261
zmap           74613
nmap_connect  45882
hping_syn     43750
unicorn_syn   43039
nmap_syn      40642
unicorn_conn  39170
masscan       21138
nmap_ack      20497
nmap_window   18851
nmap_null     12511
nmap_xmas     12505
nmap_fin      12504
hping_ack     11344
nmap_maimon   10493
unicorn_ack   8494
hping_xmas    7344
hping_null    7344
hping_fin     7344
unicorn_null  4690
unicorn_xmas  4466
unicorn_fxmas 4444
unicorn_fin   4438
Name: count, dtype: int64
```

## Features that are irrelevant to my objective

Removed layer-2 related features as below:

- frame\_info.time
- frame\_info.encap\_type
- frame\_info.time\_epoch
- frame\_info.number
- frame\_info.len
- frame\_info.cap\_len
- eth.type

Removed redundant or constant features as below:

- ip.version - only IPV4 is taken now
- ip.proto - only TCP records/packets are taken now
- ip.src
- ip.dst
- ip.flags
- tcp.flags

```
In [16]: full_data.drop(columns=['frame_info.time', 'frame_info.encap_type', 'frame_info.
frame_info.len', 'frame_info.cap_len', 'eth.type', 'ip.
ip.version', 'ip.proto', 'tcp.flags'], axis=1, inplace=
```

## Exploratory Data Analysis

In [17]: `full_data.info()`

```
<class 'pandas.core.frame.DataFrame'>
Index: 783764 entries, 1 to 455502
Data columns (total 29 columns):
#   Column                Non-Null Count  Dtype
---  -
0   ip.hdr_len             783764 non-null float64
1   ip.tos                 783764 non-null float64
2   ip.id                  783764 non-null int64
3   ip.flags.rb            783764 non-null float64
4   ip.flags.df            783764 non-null float64
5   ip.flags.mf            783764 non-null float64
6   ip.frag_offset         783764 non-null float64
7   ip.ttl                 783764 non-null float64
8   ip.checksum            783764 non-null int64
9   ip.len                 783764 non-null float64
10  ip.dsfield              783764 non-null int64
11  tcp.srcport            783764 non-null float64
12  tcp.dstport            783764 non-null float64
13  tcp.seq                783764 non-null float64
14  tcp.ack                783764 non-null float64
15  tcp.len                783764 non-null float64
16  tcp.hdr_len            783764 non-null float64
17  tcp.flags.fin          783764 non-null float64
18  tcp.flags.syn          783764 non-null float64
19  tcp.flags.reset        783764 non-null float64
20  tcp.flags.push         783764 non-null float64
21  tcp.flags.ack          783764 non-null float64
22  tcp.flags.urg          783764 non-null float64
23  tcp.flags.cwr          783764 non-null float64
24  tcp.window_size        783764 non-null float64
25  tcp.checksum           783764 non-null int64
26  tcp.urgent_pointer     783764 non-null float64
27  tcp.options.mss_val    783764 non-null float64
28  label                  783764 non-null object
dtypes: float64(24), int64(4), object(1)
memory usage: 179.4+ MB
```

In [18]: `full_data.describe()`

Out[18]:

	ip.hdr_len	ip.tos	ip.id	ip.flags.rb	ip.flags.df	ip.flags.mf	ip
<b>count</b>	783764.0	783764.0	783764.000000	783764.0	783764.000000	783764.000000	
<b>mean</b>	20.0	0.0	30203.740695	0.0	0.645392	0.000003	
<b>std</b>	0.0	0.0	20375.832825	0.0	0.478395	0.001597	
<b>min</b>	20.0	0.0	0.000000	0.0	0.000000	0.000000	
<b>25%</b>	20.0	0.0	11779.000000	0.0	0.000000	0.000000	
<b>50%</b>	20.0	0.0	29827.000000	0.0	1.000000	0.000000	
<b>75%</b>	20.0	0.0	48693.000000	0.0	1.000000	0.000000	
<b>max</b>	20.0	0.0	65535.000000	0.0	1.000000	1.000000	

## Removed columns with zero variance as they dont support learning

```
In [19]: full_data.drop(columns=['label']).var() == 0
```

```
Out[19]: ip.hdr_len      True
ip.tos        True
ip.id         False
ip.flags.rb   True
ip.flags.df   False
ip.flags.mf   False
ip.frag_offset True
ip.ttl        False
ip.checksum   False
ip.len        False
ip.dsfield    False
tcp.srcport   False
tcp.dstport   False
tcp.seq       False
tcp.ack       False
tcp.len       False
tcp.hdr_len   False
tcp.flags.fin  False
tcp.flags.syn  False
tcp.flags.reset False
tcp.flags.push False
tcp.flags.ack  False
tcp.flags.urg  False
tcp.flags.cwr  False
tcp.window_size False
tcp.checksum   False
tcp.urgent_pointer False
tcp.options.mss_val False
dtype: bool
```


```
In [20]: full_data.drop(columns=['ip.hdr_len', 'ip.tos', 'ip.flags.rb',
                                'ip.frag_offset', 'ip.flags.mf'], axis=1, inplace=True)
```

## Linear Correlation greater than 0

```
In [21]: corr_data = full_data.drop(columns=['label']).corr().abs().round(2)

# get only the features with positive correlation
corr_data = corr_data[corr_data>0]

plt.figure(figsize = (20,16))
sns.heatmap(corr_data, xticklabels=corr_data.columns, yticklabels=corr_data.colu
```



binary classification (attack or bonafide)

```
In [22]: full_data['label'].value_counts()
```

```
Out[22]: label
         bonafide      328261
         zmap          74613
         nmap_connect  45882
         hping_syn     43750
         unicorn_syn   43039
         nmap_syn      40642
         unicorn_conn  39170
         masscan       21138
         nmap_ack      20497
         nmap_window   18851
         nmap_null     12511
         nmap_xmas     12505
         nmap_fin      12504
         hping_ack     11344
         nmap_maimon   10493
         unicorn_ack   8494
         hping_xmas    7344
         hping_null    7344
         hping_fin     7344
         unicorn_null  4690
         unicorn_xmas  4466
         unicorn_fxmas 4444
         unicorn_fin   4438
         Name: count, dtype: int64
```

## label encode the label column

replace bonafide labels as 0

replace attack labels as 1

```
In [23]: full_data.label[full_data.label == "bonafide"] = 0
         full_data.label[full_data.label != 0] = 1
         full_data['label'].value_counts()
```

```
Out[23]: label
         1    455503
         0    328261
         Name: count, dtype: int64
```

```
In [24]: full_data.shape
```

```
Out[24]: (783764, 24)
```

## Removal of some columns either random or might hinder learning

- acknowledgement and checksums are random
- it is known tcp.dstport promotes learning the lab architecture where dataset was collected which is UNINTENDED

```
In [25]: full_data.drop(columns=["ip.checksum", "tcp.checksum",
                                "tcp.ack", "tcp.dstport"], axis=1, inplace=True)
```



```
In [26]: full_data.head(5)
full_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 783764 entries, 1 to 455502
Data columns (total 20 columns):
#   Column                Non-Null Count  Dtype
---  -
0   ip.id                  783764 non-null  int64
1   ip.flags.df            783764 non-null  float64
2   ip.ttl                 783764 non-null  float64
3   ip.len                 783764 non-null  float64
4   ip.dsfield             783764 non-null  int64
5   tcp.srcport            783764 non-null  float64
6   tcp.seq                783764 non-null  float64
7   tcp.len                783764 non-null  float64
8   tcp.hdr_len            783764 non-null  float64
9   tcp.flags.fin          783764 non-null  float64
10  tcp.flags.syn          783764 non-null  float64
11  tcp.flags.reset        783764 non-null  float64
12  tcp.flags.push         783764 non-null  float64
13  tcp.flags.ack           783764 non-null  float64
14  tcp.flags.urg           783764 non-null  float64
15  tcp.flags.cwr           783764 non-null  float64
16  tcp.window_size         783764 non-null  float64
17  tcp.urgent_pointer      783764 non-null  float64
18  tcp.options.mss_val     783764 non-null  float64
19  label                  783764 non-null  object
dtypes: float64(17), int64(2), object(1)
memory usage: 125.6+ MB
```

## Drop duplicates records

```
In [27]: full_data.drop_duplicates(inplace=True, ignore_index=True)
```

```
In [28]: full_data['label'].value_counts()
```

```
Out[28]: label
1      362150
0      315729
Name: count, dtype: int64

=====
```

## save preprocessed dataset into CSV file

```
In [29]: full_data.to_csv("full_data_preproces_main.csv", index=False)
```

## preparation of dataset

```
In [30]: full_data = full_data.fillna(0)
X = full_data.drop(columns = ["label"])
y = full_data.label

print(X.shape, y.shape)
```

```
(677879, 19) (677879,)
```

```
In [31]: X = X.astype(int)
X.head()
X.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 677879 entries, 0 to 677878
Data columns (total 19 columns):
#   Column                Non-Null Count  Dtype
---  -
0   ip.id                  677879 non-null  int64
1   ip.flags.df            677879 non-null  int64
2   ip.ttl                 677879 non-null  int64
3   ip.len                 677879 non-null  int64
4   ip.dsfield             677879 non-null  int64
5   tcp.srcport            677879 non-null  int64
6   tcp.seq                677879 non-null  int64
7   tcp.len                677879 non-null  int64
8   tcp.hdr_len            677879 non-null  int64
9   tcp.flags.fin          677879 non-null  int64
10  tcp.flags.syn          677879 non-null  int64
11  tcp.flags.reset        677879 non-null  int64
12  tcp.flags.push         677879 non-null  int64
13  tcp.flags.ack           677879 non-null  int64
14  tcp.flags.urg           677879 non-null  int64
15  tcp.flags.cwr           677879 non-null  int64
16  tcp.window_size        677879 non-null  int64
17  tcp.urgent_pointer     677879 non-null  int64
18  tcp.options.mss_val    677879 non-null  int64
dtypes: int64(19)
memory usage: 98.3 MB
```

## Grid Search for Machine Learning Algorithms

### ML model generation

```

In [34]: from sklearn.tree import DecisionTreeClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.linear_model import LogisticRegression
from xgboost import XGBClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import LinearSVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn.neural_network import MLPClassifier
from sklearn.calibration import CalibratedClassifierCV

from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import StratifiedKFold, GridSearchCV
from sklearn.metrics import make_scorer, f1_score, roc_auc_score, accuracy_score

algorithms = {
    "SVM" : (LinearSVC(random_state=17), {}),
    "KNN" : (KNeighborsClassifier(n_jobs=-1), {
        "n_neighbors" : [1, 3, 5]
    }),
    "MLP" : (MLPClassifier(random_state=17), {
        "hidden_layer_sizes" : (10, 10),
    }),
    "NB" : (GaussianNB(), {}),
    "XGB" : (XGBClassifier(random_state=17, n_jobs=-1), {}),
    "LR" : (LogisticRegression(random_state=17, n_jobs=-1), {}),

    "RF" : (RandomForestClassifier(random_state=17, n_jobs=-1), {
        "n_estimators" : [10, 15, 20],
        "criterion" : ("gini", "entropy"),
        "max_depth" : [5, 10],
        "class_weight": (None, "balanced", "balanced_subsample")
    }),
    "DT" : (DecisionTreeClassifier(random_state=17), {
        "criterion": ("gini", "entropy"),
        "max_depth": [5, 10, 15],
        "class_weight": (None, "balanced")
    }),
}

```

## Best parameters estimation by f1-score

```

In [35]: # Train, Test
k_fold_cv = StratifiedKFold(n_splits=10, shuffle=True, random_state=17)
# Validation
grid_search_k_fold = StratifiedKFold(n_splits=3, shuffle=True, random_state=17)
# Performance metric
# f1-score; it can be considered roc_auc score for binary classification (attack

performance_scores = {}
best_parameters = {}
# placeholder for algos and their scores
for algo in algorithms.keys():
    performance_scores[algo] = { 'actual': [], 'predicted': [] }

for algo, (model, def_params) in algorithms.items():
    print(algo)
    # inner loop runs k times eql to num of folds
    for train_fold, test_fold in k_fold_cv.split(X, y):
        # scale the train fold set
        scaler = StandardScaler()
        X_train = scaler.fit_transform(X.iloc[train_fold])
        y_train = y.iloc[train_fold]

        # fit the parameter estimator
        estimator = GridSearchCV(model, def_params, cv=grid_search_k_fold, scoring=m
        estimator.fit(X_train, y_train)
        # save the best hyperparameters
        best_parameters[algo] = estimator.best_params_

        # scale the test fold set
        X_test = scaler.transform(X.iloc[test_fold])
        y_test = y.iloc[test_fold]

        # predict by the best candidate model
        y_pred = estimator.predict(X_test)

        # saving the y_test and y_pred for later evaluation
        # FOR ALL TEST FOLDS COMBINED
        performance_scores[algo]['actual'].extend(y_test)
        performance_scores[algo]['predicted'].extend(y_pred)

```

SVM  
 KNN  
 MLP  
 NB  
 XGB  
 LR  
 RF  
 DT

## Grid search results of best model hyperparameters by f1-score

```

In [37]: best_parameters

```

```
Out[37]: {'SVM': {},
          'KNN': {'n_neighbors': 1},
          'MLP': {'hidden_layer_sizes': 10},
          'NB': {},
          'XGB': {},
          'LR': {},
          'RF': {'class_weight': 'balanced_subsample',
                 'criterion': 'gini',
                 'max_depth': 10,
                 'n_estimators': 20},
          'DT': {'class_weight': None, 'criterion': 'entropy', 'max_depth': 15}}
```

## ROC\_AUC evaluation for best hyperparameter selection

best set of parameter estimation by roc\_auc

```
In [38]: algorithms2 = {
          "SVM" : (CalibratedClassifierCV(LinearSVC(random_state=17), n_jobs=-1), {}),
          "MLP" : (MLPClassifier(random_state=17), {
                    "hidden_layer_sizes" : (10, 10),
                  }),
          "KNN" : (KNeighborsClassifier(n_jobs=-1), {
                    "n_neighbors" : [1, 3, 5]
                  }),
          "XGB" : (XGBClassifier(random_state=17, n_jobs=-1), {}),
          "NB" : (GaussianNB(), {}),
          "LR" : (LogisticRegression(random_state=17, n_jobs=-1), {}),
          "RF" : (RandomForestClassifier(random_state=17, n_jobs=-1), {
                    "n_estimators" : [10, 15, 20],
                    "criterion" : ("gini", "entropy"),
                    "max_depth": [5, 10],
                    "class_weight": (None, "balanced", "balanced_subsample")
                  }),
          "DT" : (DecisionTreeClassifier(random_state=17), {
                    "criterion": ("gini", "entropy"),
                    "max_depth": [5, 10, 15],
                    "class_weight": (None, "balanced")
                  }),
        }
```

```

In [39]: # Train, Test
k_fold_cv = StratifiedKFold(n_splits=10, shuffle=True, random_state=17)
# Validation
grid_search_k_fold = StratifiedKFold(n_splits=3, shuffle=True, random_state=17)
# Performance metric
# roc_auc_score

performance_scores2 = {}
best_parameters2 = {}
for algo in algorithms2.keys():
    performance_scores2[algo] = {'actual': [], 'predicted': []}

for algo, (model, def_params) in algorithms2.items():
    print(algo)
    # inner loop runs k times eql to num of folds
    for train_fold, test_fold in k_fold_cv.split(X, y):

        # scale the train fold set
        scaler = StandardScaler()
        X_train = scaler.fit_transform(X.iloc[train_fold])
        y_train = y.iloc[train_fold]

        # fit the parameter estimator
        estimator = GridSearchCV(model, def_params, cv=grid_search_k_fold, scoring=m
estimator.fit(X_train, y_train)

        # save the best hyperparameters
        best_parameters2[algo] = estimator.best_params_

        # scale the test fold set
        X_test = scaler.transform(X.iloc[test_fold])
        y_test = y.iloc[test_fold]

        # predict by the best candidate model
        y_pred = estimator.predict_proba(X_test).transpose()[1]

        # saving y_test and y_pred for later evaluation
        performance_scores2[algo]['actual'].extend(y_test)
        performance_scores2[algo]['predicted'].extend(y_pred)

```

SVM  
 MLP  
 KNN  
 XGB  
 NB  
 LR  
 RF  
 DT

```

In [41]: best_parameters2

```

```
Out[41]: {'SVM': {},
          'MLP': {'hidden_layer_sizes': 10},
          'KNN': {'n_neighbors': 1},
          'XGB': {},
          'NB': {},
          'LR': {},
          'RF': {'class_weight': 'balanced_subsample',
                 'criterion': 'gini',
                 'max_depth': 10,
                 'n_estimators': 20},
          'DT': {'class_weight': None, 'criterion': 'entropy', 'max_depth': 15}}
```

## Evaluation metric scores of the best candidate model for each Algorithm on all TEST FOLDS combined

classification metrics scores for best set of parameters (candidate model) on ALL TEST FOLDS combined

```
In [ ]: from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score

metrics = [accuracy_score, precision_score, recall_score, f1_score, roc_auc_score]
score_names = ['Accuracy', 'Precision', 'Recall', 'F1', 'ROC']
scores = {}
for name in score_names:
    scores[name] = {}

# calculate the metrics scores
for i in range(len(metrics)):
    temp = {}
    for algo in algorithms.keys():
        if metrics[i] is metrics[-1]: # include the roc scores
            temp.update({algo: metrics[i](performance_scores2[algo]['actual'], performance_scores2[algo]['predicted'])})
        else: # include rest of the scores
            temp.update({algo: metrics[i](performance_scores[algo]['actual'], performance_scores[algo]['predicted'])})
    scores[score_names[i]] = dict(sorted(temp.items(), key=lambda item: item[1], reverse=True))
    del temp # release memory
```

```
In [254... for key, value in scores.items():
            print(key)
            for k, v in value.items():
                print(k, v)
            print("\n")
```

## Accuracy

DT 0.9999365668504261  
XGB 0.9998731337008522  
KNN 0.9991591419707647  
RF 0.9989363883524935  
MLP 0.9987696919361715  
LR 0.9084556388381997  
SVM 0.9016668166442684  
NB 0.6913977273230179

## Precision

DT 0.9998950806636424  
XGB 0.9997984594572209  
RF 0.9996074235650917  
KNN 0.9989072003355668  
MLP 0.9987135884808198  
LR 0.8835125906084182  
SVM 0.8668393113414143  
NB 0.6338608243001745

## Recall

DT 0.9999861935662019  
NB 0.9999668645588844  
XGB 0.9999641032721248  
KNN 0.9995195361038244  
MLP 0.9989838464724562  
RF 0.9984012149661743  
SVM 0.9640287173823001  
LR 0.9544912329145382

## F1

DT 0.9999406350394087  
XGB 0.9998812745043901  
KNN 0.9992132744071264  
RF 0.9990039551683122  
MLP 0.998848699195746  
LR 0.9176314119836367  
SVM 0.9128544235616496  
NB 0.7758955554055776

## ROC

XGB 0.9999988265040711  
RF 0.9999860555451553  
DT 0.9999408173859827  
MLP 0.9999251679210259  
KNN 0.9991326479584143  
LR 0.9696534960046947  
SVM 0.9617430380478496  
NB 0.924145986912221



```
In [226... for i in range(len(score_names)):

    print(score_names[i])
    score_df = pd.DataFrame(data=scores[score_names[i]], columns=scores[score_na
    score_df.to_csv(f"scores/{score_names[i]}_scores.csv", index=True)
```

Accuracy  
Precision  
Recall  
F1  
ROC

```
In [216... score_df
```

```
Out[216]:
```

	XGB	RF	DT	MLP	KNN	LR	SVM	NB
<b>ROC</b>	0.999999	0.999986	0.999941	0.999925	0.999133	0.969653	0.961743	0.924146

```
In [ ]:
```

## plotting the ROC/AUC curve

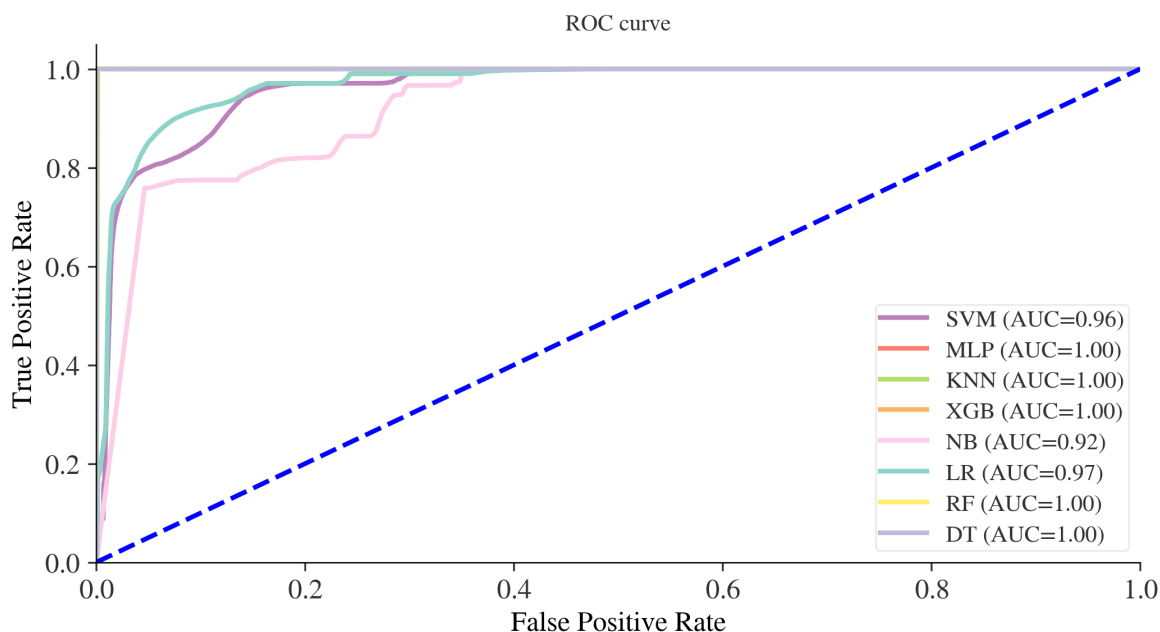
```
In [265... import os
from sklearn.metrics import auc

plt.figure(figsize=(8,4))

for algo, score in performance_scores2.items():
    fpr, tpr, threshold = roc_curve(score['actual'], score['predicted'])
    auc_score = auc(fpr,tpr)
    plt.plot(fpr, tpr, label="{0} (AUC={:.2f})".format(algo, auc_score))

plt.plot([0,1], [0,1], color='blue', linestyle='--', linewidth=2)
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC curve', fontsize=10, loc='center')
plt.legend(loc="lower right")

plt.savefig('figures/roc_curve_models.jpg', dpi=300, bbox_inches="tight")
```



## plotting the scores of evaluation metrics

```
In [236... def barplot( classifiers, scores, title, xlabel, ylabel):
    # Create a bar graph
    plt.figure(figsize=(12, 6)) # Set the size of the figure
    bars = plt.bar(classifiers, scores, color=['#1f77b4', '#ff7f0e', '#2ca02c',

    # Adding title and Labels
    plt.title(title, fontsize=16)
    plt.xlabel(xlabel, fontsize=14)
    plt.ylabel(ylabel, fontsize=14)
    plt.ylim(0, 1) # Set the y-axis range from 0 to 1

    # Adding the data Labels on top of the bars
    for bar in bars:
        yval = bar.get_height()
        label_pos = yval - yval/10
        plt.text(bar.get_x() + bar.get_width()/2, label_pos, round(yval, 2), va=

    # Improve layout and show the plot
    plt.xticks(rotation=45, ha='right') # Rotate x-axis labels for better reada
    # plt.tight_layout()
    plt.grid(axis='y', linestyle='--', alpha=0.9) # Add a grid for y-axis
    plt.savefig('figures/'+ylabel, dpi=300, bbox_inches="tight")
    plt.show()
```


## accuracy scores of classifiers

```
In [229... accuracy_df = pd.read_csv('scores/Accuracy_scores.csv', index_col=0 )
accuracy_df
```

```
Out[229]:
```

	DT	XGB	KNN	RF	MLP	LR	SVM	NB
<b>Accuracy</b>	0.999937	0.999873	0.999159	0.998936	0.99877	0.908456	0.901667	0.691398

```
In [245... barplot(classifiers=accuracy_df.columns, scores=accuracy_df.iloc[0], title="Accu  
xlabel="Classifiers", ylabel='Accuracy Score')
```




## precision scores of classifiers

```
In [230... precision_df = pd.read_csv('scores/Precision_scores.csv', index_col=0 )  
precision_df
```

```
Out[230]:
```

	DT	XGB	RF	KNN	MLP	LR	SVM	NB
<b>Precision</b>	0.999895	0.999798	0.999607	0.998907	0.998714	0.883513	0.866839	0.633861

```
In [246... barplot(classifiers=precision_df.columns, scores=precision_df.iloc[0], title="Pr  
xlabel="Classifiers", ylabel='Precision Score')
```



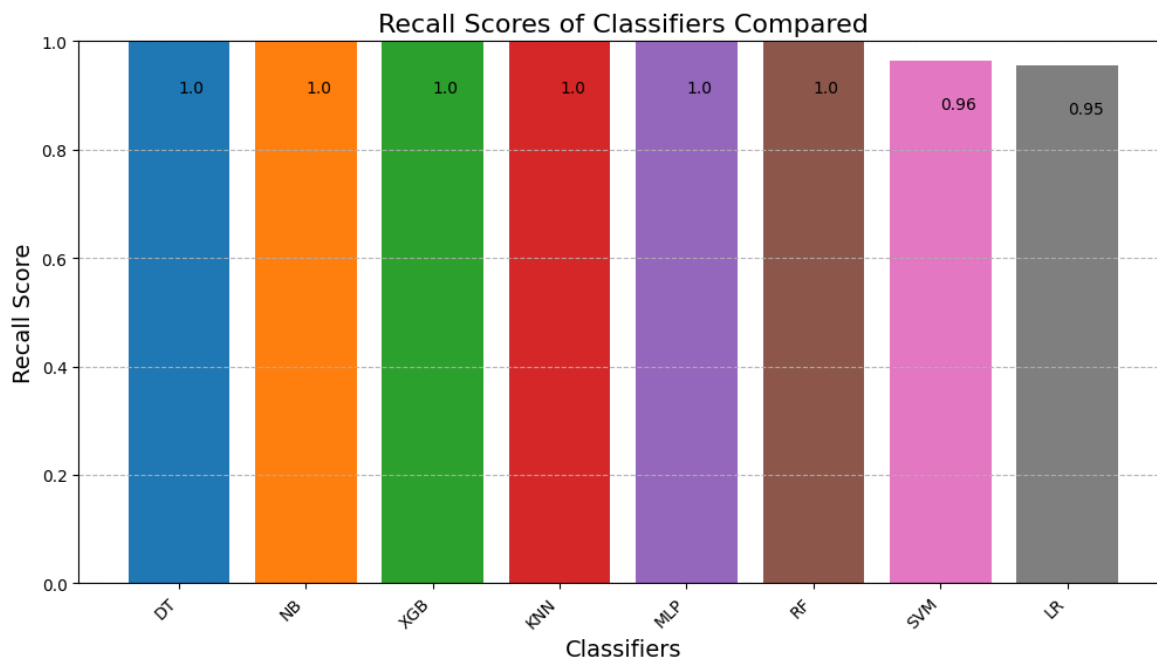
## recall scores of classifiers

```
In [233...] recall_df = pd.read_csv('scores/Recall_scores.csv', index_col=0 )  
recall_df
```

```
Out[233]:
```

	DT	NB	XGB	KNN	MLP	RF	SVM	LR
Recall	0.999986	0.999967	0.999964	0.99952	0.998984	0.998401	0.964029	0.954491

```
In [247...] barplot(classifiers=recall_df.columns, scores=recall_df.iloc[0], title="Recall S  
xlabel="Classifiers", ylabel='Recall Score')
```




## recall scores of classifier

```
In [234...] f1_df = pd.read_csv('scores/F1_scores.csv', index_col=0)  
f1_df
```

```
Out[234]:
```

	DT	XGB	KNN	RF	MLP	LR	SVM	NB
F1	0.999941	0.999881	0.999213	0.999004	0.998849	0.917631	0.912854	0.775896

```
In [248...] barplot(classifiers=f1_df.columns, scores=f1_df.iloc[0], title="F1 Scores of Cla  
xlabel="Classifiers", ylabel='F1 Score')
```



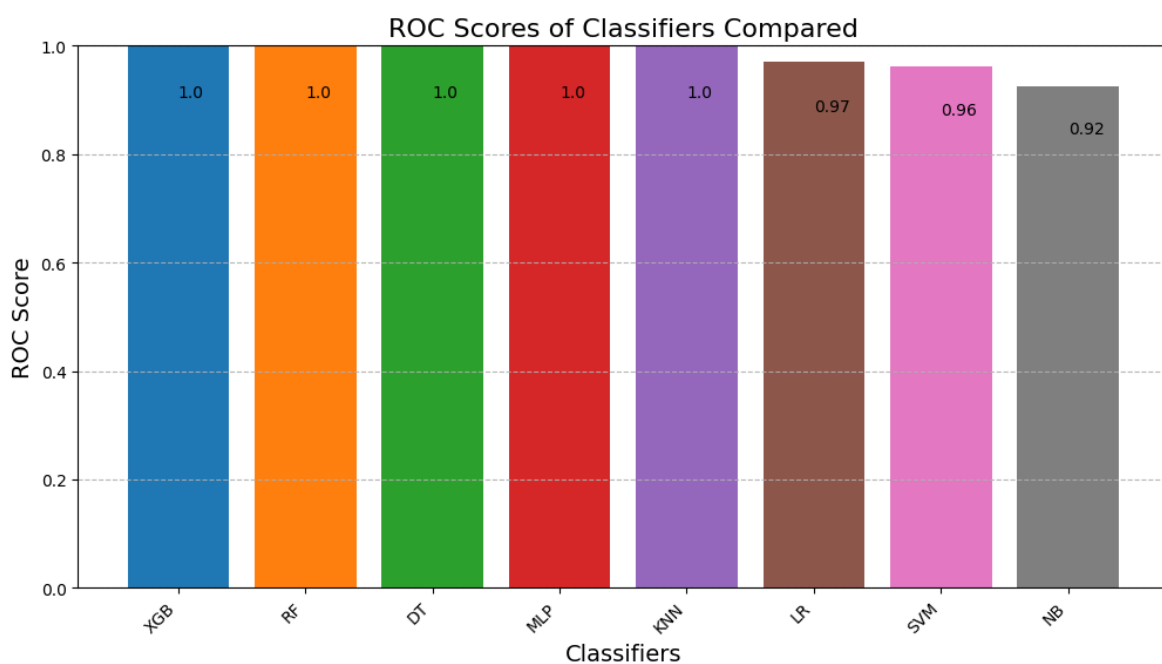
## roc scores for classifiers

```
In [235...] roc_df = pd.read_csv('scores/ROC_scores.csv', index_col=0)
roc_df
```

```
Out[235]:
```

	XGB	RF	DT	MLP	KNN	LR	SVM	NB
<b>ROC</b>	0.999999	0.999986	0.999941	0.999925	0.999133	0.969653	0.961743	0.924146

```
In [249...] barplot(classifiers=roc_df.columns, scores=roc_df.iloc[0], title="ROC Scores of
xlabel="Classifiers", ylabel='ROC Score')
```



```
In [ ]:
```